

# Deep Learning Project: Synthesis

Generating images by training Generative Adversarial Networks (GANs)

---

Johannes Loevenich

January 29, 2019

University of Bonn

# Table of contents

1. Motivation
2. Training experiences
3. Theory
4. Experiment 1: Classification of Food-101
5. Experiment 2: LSUN conference-room
6. Experiment 3: Vector arithmetics for celebA

# Motivation

---

## Motivation

- Understanding Generative Adversarial models especially DCGAN's.

## Motivation

- Understanding Generative Adversarial models especially DCGAN's.
- Can we measure the quality of the **discriminator** by removing the real/fake classifier and feeding the convolutional features into a new classifier? This would show that the model learned general, useful features.

# GENERATIVE ADVERSARIAL NETWORKS

## Motivation

- Understanding Generative Adversarial models especially DCGAN's.
- Can we measure the quality of the **discriminator** by removing the real/fake classifier and feeding the convolutional features into a new classifier? This would show that the model learned general, useful features.
- Is it possible to show that there are dedicated parts of the **generator** that control properties of its output? In other words, do we reach some vector arithmetics on the input vector noise.

## Training experiences

---

## Motivation

- Training a DCGAN is some kind of an art.



## Motivation

- Training a DCGAN is some kind of an art.
- While training the DCGAN on IMAGENET-1K the **discriminator** always got too strong.

## Motivation

- Training a DCGAN is some kind of an art.
- While training the DCGAN on IMAGENET-1K the **discriminator** always got too strong.
- We had to interrupt training periodically and then just trained the **generator** for a few epochs.

# GENERATIVE ADVERSARIAL NETWORKS

## Motivation

- Training a DCGAN is some kind of an art.
- While training the DCGAN on IMAGENET-1K the **discriminator** always got too strong.
- We had to interrupt training periodically and then just trained the **generator** for a few epochs.
- With our GPU time it was not possible to reach nearly a good classification accuracy as they did in the papers.

# GENERATIVE ADVERSARIAL NETWORKS

## Motivation

- Training a DCGAN is some kind of an art.
- While training the DCGAN on IMAGENET-1K the **discriminator** always got too strong.
- We had to interrupt training periodically and then just trained the **generator** for a few epochs.
- With our GPU time it was not possible to reach nearly a good classification accuracy as they did in the papers.
- Training on celebA worked well and we got some nice results for the vector arithmetics.

# Theory

---

# GENERATIVE ADVERSARIAL NETWORKS

## What is a GAN?

- GANs are a framework for teaching a DL model to capture the training data's distribution so we can generate new data from that same distribution.

## What is a GAN?

- GANs are a framework for teaching a DL model to capture the training data's distribution so we can generate new data from that same distribution.
- They are made of two distinct models, a **generator** and a **discriminator**.

# GENERATIVE ADVERSARIAL NETWORKS

## What is a GAN?

- GANs are a framework for teaching a DL model to capture the training data's distribution so we can generate new data from that same distribution.
- They are made of two distinct models, a **generator** and a **discriminator**.
- The job of the **generator** is to spawn fake images that look like the training images.



# GENERATIVE ADVERSARIAL NETWORKS

## What is a GAN?

- GANs are a framework for teaching a DL model to capture the training data's distribution so we can generate new data from that same distribution.
- They are made of two distinct models, a **generator** and a **discriminator**.
- The job of the **generator** is to spawn fake images that look like the training images.
- The job of the **discriminator** is to look at an image and output whether or not it is a real training image or a fake image from the generator.

## Discriminator Notation

- Let  $x$  be data representing an image.  $D(x)$  is the discriminator network which outputs the (scalar) probability that  $x$  came from training data rather than the generator.

## Discriminator Notation

- Let  $x$  be data representing an image.  $D(x)$  is the discriminator network which outputs the (scalar) probability that  $x$  came from training data rather than the generator.
- Here, since we are dealing with images the input to  $D(x)$  is an image of HWC size  $3 \times 64 \times 64$ .

## Discriminator Notation

- Let  $x$  be data representing an image.  $D(x)$  is the discriminator network which outputs the (scalar) probability that  $x$  came from training data rather than the generator.
- Here, since we are dealing with images the input to  $D(x)$  is an image of HWC size  $3 \times 64 \times 64$ .
- Intuitively,  $D(x)$  should be HIGH when  $x$  comes from training data and LOW when  $x$  comes from the generator.

## Generator Notation

- For the generators notation, let  $z$  be a latent space vector sampled from a standard normal distribution.

# GENERATIVE ADVERSARIAL NETWORKS

## Generator Notation

- For the generators notation, let  $z$  be a latent space vector sampled from a standard normal distribution.
- $G(z)$  represents the generator function which maps the latent vector  $z$  to data-space.

# GENERATIVE ADVERSARIAL NETWORKS

## Generator Notation

- For the generators notation, let  $z$  be a latent space vector sampled from a standard normal distribution.
- $G(z)$  represents the generator function which maps the latent vector  $z$  to data-space.
- The goal of  $G$  is to estimate the distribution that the training data comes from  $p_{data}$  so it can generate fake samples from that estimated distribution  $p_g$ .

## Generator Notation

- For the generators notation, let  $z$  be a latent space vector sampled from a standard normal distribution.
- $G(z)$  represents the generator function which maps the latent vector  $z$  to data-space.
- The goal of  $G$  is to estimate the distribution that the training data comes from  $p_{data}$  so it can generate fake samples from that estimated distribution  $p_g$ .
- So,  $D(G(z))$  is the probability (scalar) that the output of the generator  $G$  is a real image.



# GENERATIVE ADVERSARIAL NETWORKS

## Training by playing a MinMax game

- As described in Goodfellow's paper,  $D$  and  $G$  play a **MinMax game** in which  $D$  tries to maximize the probability it correctly classifies reals and fakes ( $\log D(x)$ ), and  $G$  tries to minimize the probability that  $D$  will predict its outputs are fake ( $\log(1 - D(G(x)))$ ).

# GENERATIVE ADVERSARIAL NETWORKS

## Training by playing a MinMax game

- As described in Goodfellow's paper,  $D$  and  $G$  play a **MinMax game** in which  $D$  tries to maximize the probability it correctly classifies reals and fakes ( $\log D(x)$ ), and  $G$  tries to minimize the probability that  $D$  will predict its outputs are fake ( $\log(1 - D(G(x)))$ ).
- The GAN **loss function** is:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(x)))]$$

# GENERATIVE ADVERSARIAL NETWORKS

## Training by playing a MinMax game

- As described in Goodfellow's paper,  $D$  and  $G$  play a **MinMax game** in which  $D$  tries to maximize the probability it correctly classifies reals and fakes ( $\log D(x)$ ), and  $G$  tries to minimize the probability that  $D$  will predict its outputs are fake ( $\log(1 - D(G(x)))$ ).
- The GAN **loss function** is:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(x)))]$$

- In theory, the solution to this **MinMax game** is where  $p_g = p_{data}$  and the discriminator guesses randomly if the inputs are real or fake.

# GENERATIVE ADVERSARIAL NETWORKS

## What is a DCGAN?

- A **DCGAN** is a direct extension of the **GAN** described above, except that it explicitly uses convolutional and convolutional-transpose layers in the discriminator and generator, respectively.

# GENERATIVE ADVERSARIAL NETWORKS

## What is a DCGAN?

- A **DCGAN** is a direct extension of the **GAN** described above, except that it explicitly uses convolutional and convolutional-transpose layers in the discriminator and generator, respectively.
- The **discriminator** is made up of strided convolution layers, batch norm layers, and LeakyReLU activations.

# GENERATIVE ADVERSARIAL NETWORKS

## What is a DCGAN?

- A **DCGAN** is a direct extension of the **GAN** described above, except that it explicitly uses convolutional and convolutional-transpose layers in the discriminator and generator, respectively.
- The **discriminator** is made up of strided convolution layers, batch norm layers, and LeakyReLU activations.
- The input is a 3x64x64 input image and the output is a scalar probability that the input is from the real data distribution.

# GENERATIVE ADVERSARIAL NETWORKS

## What is a DCGAN?

- A **DCGAN** is a direct extension of the **GAN** described above, except that it explicitly uses convolutional and convolutional-transpose layers in the discriminator and generator, respectively.
- The **discriminator** is made up of strided convolution layers, batch norm layers, and LeakyReLU activations.
- The input is a  $3 \times 64 \times 64$  input image and the output is a scalar probability that the input is from the real data distribution.
- The **generator** is comprised of convolutional-transpose layers, batch norm layers, and ReLU activations.

# GENERATIVE ADVERSARIAL NETWORKS

## What is a DCGAN?

- A **DCGAN** is a direct extension of the **GAN** described above, except that it explicitly uses convolutional and convolutional-transpose layers in the discriminator and generator, respectively.
- The **discriminator** is made up of strided convolution layers, batch norm layers, and LeakyReLU activations.
- The input is a  $3 \times 64 \times 64$  input image and the output is a scalar probability that the input is from the real data distribution.
- The **generator** is comprised of convolutional-transpose layers, batch norm layers, and ReLU activations.
- The input is a latent vector,  $z$ , that is drawn from a standard normal distribution and the output is a  $3 \times 64 \times 64$  RGB image.



# STRUCTURE OF THE DCGAN NETWORK




Figure 1: DCGAN network

# **Experiment 1: Classification of Food-101**

---

# CLASSIFICATION ON FOOD-101

## The Data:

- IMAGENET-1K: pictures of 1000 different objects
- food-101: 101.000 images in 101 different food categories

## The Experiment:

- Train a DCGAN on the IMAGENET-1K dataset. Use the **discriminator** as feature extractor for food-101.
- Use the output of all the convolutional layers of the **discriminator** as feature extractor and train a linear SVM on food-101.

# RESULTS CLASSIFICATION ON FOOD-101

## Overall accuracy

- We tested the **discriminator** as feature extractor on the whole food-101 dataset (101.000 pictures in 101 categories) and on 10 random categories.

# RESULTS CLASSIFICATION ON FOOD-101

## Overall accuracy

- We tested the **discriminator** as feature extractor on the whole food-101 dataset (101.000 pictures in 101 categories) and on 10 random categories.
- To measure our classification results we used the models overall **accuracy** and **precision** and **recall** on each of datasets classes.

# RESULTS CLASSIFICATION ON FOOD-101

## Overall accuracy

- We tested the **discriminator** as feature extractor on the whole food-101 dataset (101.000 pictures in 101 categories) and on 10 random categories.
- To measure our classification results we used the models overall **accuracy** and **precision** and **recall** on each of datasets classes.
- While **recall** expresses the ability to find all relevant instances in a dataset, **precision** expresses the proportion of the data points our model says was relevant actually were relevant.

$$recall = \frac{tp}{tp + fn}$$

$$precision = \frac{tp}{tp + fp}$$

# RESULTS CLASSIFICATION ON FOOD-101


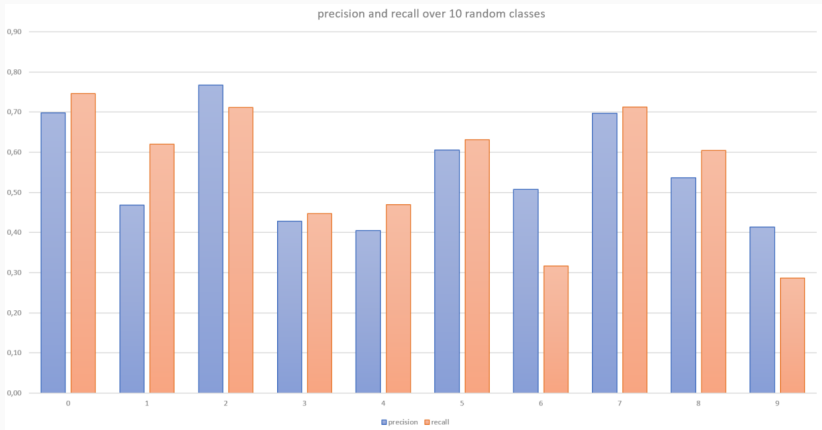


Figure 2: precision vs. recall over all 101 food categories

# RESULTS CLASSIFICATION ON FOOD-101



**Figure 3:** precision vs. recall over 10 random food categories



## **Experiment 2: LSUN conference-room**

---

## The Data:

- LSUN conference room: 229.069 images of conference-rooms


## The Experiment:

- Train a **Generator** on the LSUN conference-room dataset.

# LSUN CONFERENCE-ROOM

## Generator and Discriminator losses


- $Loss_D$  : **discriminator loss** calculated as the sum of losses for the all real and all fake batches  $\log(D(x)) + \log(D(G(z)))$ .




# LSUN CONFERENCE-ROOM

## Generator and Discriminator losses

- $Loss_D$  : **discriminator loss** calculated as the sum of losses for the all real and all fake batches  $\log(D(x)) + \log(D(G(z)))$ .
- $Loss_G$  : **generator loss** calculated as  $\log(D(G(z)))$




# LSUN CONFERENCE-ROOM REAL VS. FAKE



**Figure 5:** Real samples from the LSUN conference-room dataset

# LSUN CONFERENCE-ROOM REAL VS. FAKE



**Figure 6:** Fake samples after 25 epochs

## Experiment 3: Vector arithmetics for celebA

---

# VECTOR ARITHMETICS FOR CELEBA

## The Data:


- celebA faces dataset: 10.177 number of identities , 202.599 number of face images

## The Experiment:

- Train a **DCGAN** on the celebA faces dataset. Take two sets of 3 normal distributed sample vectors  $A$  and  $B$ . Calculate both subsets means ( $mean(A)$  and  $mean(B)$ ). Then return the **generator's** output of  $mean(a) + mean(B)$ .



# RESULTS VECTOR ARITHMETICS FOR CELEBA



**Figure 7: Means**

# RESULTS VECTOR ARITHMETICS FOR CELEBA

A\_mean + B\_mean



**Figure 8:** Sum of the means

**Questions?**